

Rackspace Fine Tuning-as-a-Service

Accelerate AI Adoption with Models Built for Your Business

Solution Overview

Rackspace Fine-tuning as a Service provides enterprise-grade fine-tuning capabilities designed to help organizations customize open-source AI models using proprietary business data, workflows, and domain-specific knowledge. Built on a Kubernetes-native architecture and optimized for both NVIDIA and AMD GPU platforms, it enables organizations to securely train and optimize models for production AI use cases without the complexity of managing AI training infrastructure internally. The platform supports modern fine-tuning methodologies, including LoRA and QLoRA, enabling efficient model customization, reducing infrastructure requirements, and accelerating training cycles.


Designed for enterprise AI adoption at scale, Rackspace FTaaS combines scalable GPU infrastructure, managed training operations, and integrated model lifecycle management into a unified AI training platform. Organizations can securely upload and manage datasets, orchestrate distributed training workloads, and operationalize fine-tuned models across private, hybrid, and air-gapped environments. With integrated observability, operational governance, and enterprise security controls, it enables organizations to accelerate AI innovation while maintaining control, performance, and operational consistency across the AI development lifecycle.

Key Capabilities

Rackspace FTaaS Inferencing supports:

- LoRA and QLoRA fine-tuning
- Dataset upload and management
- Distributed training orchestration
- Experiment tracking
- Model versioning and lifecycle management
- GPU-aware training optimization
- NVIDIA and AMD GPU support
- Kubernetes-native training infrastructure
- Dedicated and isolated training environments
- Secure model artifact management

Rackspace FTaaS Environment



Key Differentiators

- Enterprise-grade managed fine-tuning operations
- Support for open-source foundation models
- NVIDIA and AMD GPU flexibility
- Private and air-gapped deployment support
- Kubernetes-native scalability
- Integrated model lifecycle management
- Secure enterprise dataset handling
- Operational governance and observability
- Flexible deployment architectures

Rackspace Inferencing-as-a-Service

Accelerate AI Adoption with Models Built for Your Business

Target Customers and Workloads

Rackspace Fine Tuning as a Service is a managed fine-tuning service that enables organizations to customize and optimize open-source AI models using enterprise data across private, hybrid, and air-gapped environments. It is designed for organizations looking to customize foundation models using proprietary enterprise data without building and managing complex AI training infrastructure internally.

Who:

- Enterprise AI/ML teams
- Software providers
- Digital-native organizations
- Regulated industries

Company size: Enterprise to Mid Market

Focus Industries: Healthcare, BFSI

Workloads:

It is particularly well-suited for organizations developing domain-specific AI assistants, enterprise copilots, intelligent automation platforms, and fine-tuned large language models optimized for industry-specific workflows and datasets.

Business Value

Rackspace Fine Tuning as a Service helps organizations:

- Customize AI models for enterprise-specific use cases
- Improve model accuracy and response quality
- Accelerate deployment of domain-specific AI applications
- Reduce the complexity of AI training infrastructure
- Securely leverage proprietary enterprise datasets
- Operationalize fine-tuned models for production workloads
- Optimize GPU resource utilization during training operations

Persona	Focus Area	Role in Decision	Key questions they ask	Influence level
CIO	AI strategy, innovation, operational efficiency	Executive sponsor and budget owner	How quickly can we operationalize AI while maintaining governance and reducing complexity?	High
CTO	Technology architecture, AI platform strategy	Technical decision maker	Can this platform become our standard AI training and model customization layer?	High
Chief Data & AI Officer (CDAO)	Data strategy and AI outcomes	Business & Technical Champion	Can we build more accurate models using our proprietary data?	High
CFO	Investment justification and ROI	Financial Approver	How do we reduce the cost of AI development while maximizing business value?	High
CISO	Security and compliance	Approval Authority	Can we protect proprietary data and intellectual property throughout the AI lifecycle?	High

CIO Value Prop

Rackspace Fine-Tuning as a Service enables CIOs to rapidly transform open-source AI models into business-specific AI assets by securely incorporating proprietary data, workflows, and domain expertise. By eliminating the complexity of managing GPU infrastructure, training operations, and model lifecycle management, organizations can accelerate AI initiatives, improve model accuracy, and deploy production-ready AI solutions faster. The platform provides the governance, security, scalability, and deployment flexibility required to operationalize AI across private, hybrid, and air-gapped environments while maintaining full control over data and intellectual property.

CFO Value Prop

Rackspace Fine-Tuning as a Service helps organizations unlock greater value from AI investments by customizing cost-efficient open-source models instead of relying solely on expensive proprietary alternatives. Efficient fine-tuning techniques such as LoRA and QLoRA reduce GPU consumption, infrastructure costs, and training time while improving model performance for specific business use cases. With a fully managed service model, organizations avoid the capital expense and specialized staffing required to build and operate AI training environments, allowing them to scale AI initiatives with predictable costs and measurable business outcomes.

Rackspace Inferencing-as-a-Service

Accelerate AI Adoption with Models Built for Your Business

Use Case	Description	Why it's a good candidate for FTaaS	Primary owner
AI Copilots	AI assistants that help employees perform tasks, answer questions, generate content, generate information, and automate workflows. Examples include HR, finance, sales, and healthcare copilots.	Fine-tuning enables the model to learn company-specific terminology, processes, policies, and preferred responses, improving accuracy and adoption..	CIO, Head of Digital Transformation, Director of AI
Industry-Specific AI Assistants	AI solutions tailored for industries such as manufacturing, energy, retail, insurance, or telecommunications.	Industry language, regulations, workflows, and terminology differ significantly from general-purpose models, making customization essential.	Business Unit Leader, CTO, Industry Solution Owner
Customer Service Automation	AI-powered virtual agents handling customer inquiries, ticket routing, call summarization, and self-service interactions.	Fine-tuning improves response consistency, product knowledge, escalation handling, and customer experience while reducing hallucinations.	VP Customer Experience, Contact Center Director
Financial Services AI Models	AI models supporting underwriting, risk assessment, fraud detection, customer service, portfolio analysis, and compliance activities.	Financial institutions require models trained on proprietary processes, products, regulatory requirements, and risk frameworks.	Chief Data Officer, Chief Risk Officer, Head of AI
Healthcare AI Applications	Clinical documentation assistants, patient engagement tools, care management systems, and medical knowledge assistants.	Fine-tuning improves response consistency, product knowledge, escalation handling, and customer experience while reducing hallucinations.	CIO, Chief Medical Information Officer (CMIO), Chief Digital Officer
Legal and Compliance Assistants	AI solutions that analyze contracts, policies, regulations, legal documents, and compliance requirements.	Fine-tuning improves interpretation of legal language, regulatory frameworks, internal policies, and jurisdiction-specific requirements.	General Counsel, Chief Compliance Officer
Code Generation and Engineering Assistants	AI tools that assist developers with coding, testing, documentation, and software lifecycle management.	Fine-tuning allows models to learn internal coding standards, architectures, frameworks, repositories, and development practices.	CTO, VP Engineering, Head of Software Development
Internal Knowledge Assistants	AI assistants that help employees access institutional knowledge, procedures, policies, and operational information.	Fine-tuning enables the model to understand organizational vocabulary, business processes, and preferred response formats.	CIO, Knowledge Management Leader, HR Technology Leader
AI-Powered SaaS Products	Software vendors embedding AI capabilities into commercial products and customer-facing applications.	Product differentiation often requires AI models that understand specific product functionality, customer workflows, and proprietary data structures.	Chief Product Officer, VP Product Management
Retrieval-Augmented Generation (RAG) Optimization	Fine-tuning models used alongside enterprise search and RAG architectures to improve response quality and retrieval performance.	Fine-tuning improves grounding behavior, reduces hallucinations, enhances citation handling, and optimizes model responses for retrieved enterprise content.	Head of AI, Enterprise Architecture, Data Science Leader

Common Characteristics of Ideal Inferencing Workloads

The strongest candidates for Inferencing as a Service typically require:

- Real-time or near real-time responses
- Large Language Model (LLM) execution
- Variable or unpredictable demand
- High GPU consumption
- Enterprise security and compliance
- API-based integration into business applications
- Production-scale availability and reliability
- Ongoing model updates and lifecycle management



Resource Links

[Solution Hub](#)

Contacts

Offering Managers:

Amine Badaoui: Amine.Badaoui@rackspace.com

Product Marketing:

Sharon Varalli: Sharon.Varalli@rackspace.com