

# Rackspace Inferencing-as-a-Service

Accelerate Enterprise AI Adoption Without Increasing Operational Complexity

## Solution Overview

Rackspace IaaS delivers enterprise-grade AI inferencing designed for organizations that require scalable, high-performance AI services without the operational complexity of managing underlying infrastructure. Built on a Kubernetes-native architecture and optimized for both NVIDIA and AMD GPU platforms, Rackspace IaaS enables enterprises to rapidly deploy and scale large language models through secure, OpenAI-compatible APIs. The platform supports dedicated and multi-tenant inference environments, intelligent autoscaling and GPU-aware workload optimization to deliver consistent low-latency performance for production AI workloads.

Designed for modern enterprise AI adoption, RackAI combines high-performance inferencing with integrated observability, operational automation and enterprise security controls. Organizations can accelerate AI initiatives while maintaining flexibility across private cloud, on-premises and air-gapped deployments. With built-in monitoring, model orchestration and centralized operations management, RackAI provides a production-ready AI runtime platform capable of supporting business-critical AI applications at scale.


## Inference Capabilities

Rackspace IaaS Inferencing supports:

- Open-source large language models
- OpenAI-compatible APIs
- Dedicated inference environments
- Shared inference environments
- Dynamic autoscaling
- GPU-backed inferencing across NVIDIA and AMD GPU platforms
- Streaming responses
- Multi-model serving
- High-throughput inference workloads



## Rackspace IaaS Environment



## Autoscaling and Performance Optimization

Rackspace IaaS dynamically scales inference workloads based on demand and GPU resource utilization across NVIDIA and AMD GPU environments.

The platform optimizes:

- GPU allocation
- Inference throughput
- Queue management
- Model concurrency
- Resource scheduling

# Rackspace Inferencing-as-a-Service

## Accelerate Enterprise AI Adoption Without Increasing Operational Complexity

### Target Customers and Workloads

Rackspace IaaS is particularly well-suited for organizations requiring dedicated AI capacity, predictable performance, operational governance, and support for open-source AI models.

#### Who:

- Software providers,
- Digital native organizations,
- Enterprise IT teams
- AI/ML engineering teams

**Company size:** Enterprise to Mid Market

**Key Focus Industries:** Healthcare, BFSI

#### Workloads:

Rackspace IaaS is designed for organizations looking to operationalize generative AI and large language models at enterprise scale without the complexity of building and managing AI infrastructure internally. The service is ideal for enterprises developing AI-powered applications, copilots, intelligent automation solutions, and real-time inference workloads that require scalable, secure, and high-performance model serving.

### Business Value

Rackspace Inferencing as a service helps organizations:

- Accelerate deployment of AI-powered applications
- Eliminate the complexity of managing AI infrastructure
- Scale inference workloads dynamically
- Optimize GPU utilization and operational costs
- Deploy AI securely across enterprise environments
- Support production-grade AI applications with managed operations

| Persona                        | Focus Area                                                   | Role in Decision              | Key questions they ask                                                     | Influence level |
|--------------------------------|--------------------------------------------------------------|-------------------------------|----------------------------------------------------------------------------|-----------------|
| CIO                            | Business outcomes, AI strategy, innovation                   | Executive sponsor             | How does this accelerate AI adoption and support business objectives?      | Very high       |
| CTO                            | Technology architecture, AI platform strategy                | Technical decision maker      | Does this fit our AI architecture and future roadmap?                      | Very high       |
| Chief Data & AI Officer (CDAO) | AI adoption, model deployment, governance                    | Business & Technical Champion | How quickly can we operationalize AI models and scale use cases?           | Very high       |
| CFO                            | Cost control, ROI, financial governance                      | Financial Approver            | What is the expected ROI and how predictable are costs?                    | High            |
| VP of Infrastructure / Cloud   | Platform operations, infrastructure management               | Key Evaluator                 | Will this reduce operational complexity and infrastructure burden?         | High            |
| Head of AI / ML Engineering    | Model serving, inference performance, developer productivity | Technical Champion            | Can it support our LLMs, APIs, GPU requirements, and deployment workflows? | High            |
| CISO                           | Security, compliance, data protection                        | Approval Authority            | Does it meet security, regulatory, and data sovereignty requirements       | High            |

#### CIO Value Prop

Rackspace AI Inferencing as a Service enables CIOs to rapidly operationalize AI across the enterprise through a secure, scalable, and production-ready platform. Built on Kubernetes-native architecture and optimized for NVIDIA and AMD GPUs, it delivers high-performance AI inferencing with enterprise-grade security, observability, and automation. By eliminating the burden of managing AI infrastructure, operations, and model-serving environments, organizations can accelerate AI initiatives, improve time-to-value, and deploy business-critical AI applications consistently across private cloud, on-premises, and air-gapped environments while maintaining governance, performance, and control.

#### CFO Value Prop

Rackspace AI Inferencing as a Service helps CFOs realize the benefits of AI without the capital investment, operational overhead, or resource requirements of building and managing AI infrastructure internally. Through intelligent GPU utilization, automated scaling, and centralized operations management, organizations optimize infrastructure efficiency while avoiding overprovisioning and unnecessary AI spend. The result is a predictable, scalable AI operating model that reduces risk, improves asset utilization, accelerates return on AI investments, and enables the business to deploy AI initiatives with confidence.

# Rackspace Inferencing-as-a-Service

Accelerate Enterprise AI Adoption Without Increasing Operational Complexity

| Use Case                         | Description                                                                                                                                                                                       | Why it's a good candidate for IaaS                                                                                                                                                                                     | Primary owner                                                                 |
|----------------------------------|---------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|-------------------------------------------------------------------------------|
| AI Copilots                      | AI assistants that help employees perform tasks, answer questions, generate content, summarize information, and automate workflows. Examples include HR, finance, sales, and healthcare copilots. | Requires low-latency responses, scalable GPU resources, and continuous access to LLMs. Usage patterns can fluctuate significantly, making elastic inferencing highly cost-effective.                                   | Chief Digital Officer, CIO, Business Unit Leaders, AI/Innovation Teams        |
| Enterprise Chatbots              | Internal or external conversational AI systems that answer employee, customer, or partner questions using enterprise knowledge.                                                                   | Demand varies widely throughout the day. Managed inferencing provides consistent performance, high availability, and simplified model operations without requiring dedicated AI infrastructure teams.                  | CIO, Customer Experience Leaders, IT Operations, Digital Transformation Teams |
| Customer Service Automation      | AI-powered virtual agents handling customer inquiries, ticket routing, call summarization, and self-service interactions.                                                                         | High transaction volumes require scalable, production-grade inferencing with predictable latency and uptime. This allows organizations to scale during peak customer demand.                                           | VP Customer Service, Contact Center Director, COO, CIO                        |
| Document Summarization           | AI-generated summaries of contracts, claims, medical records, financial reports, research documents, and other large content repositories..                                                       | Processing often occurs in bursts and requires substantial compute resources. Inferencing as a Service enables rapid scaling without overprovisioning GPU infrastructure.                                              | Legal Operations, Healthcare Informatics, Knowledge Management Teams, CIO     |
| Enterprise Search and RAG        | Retrieval-Augmented Generation systems that combine enterprise search with LLM-generated responses using company-specific data.                                                                   | Requires constant inferencing requests against vector databases and enterprise content repositories. It ensures low-latency query performance and simplifies operational complexity.                                   | CIO, Chief Data Officer, Knowledge Management Teams, Enterprise Architecture  |
| AI-Powered SaaS Applications     | Commercial software products embedding AI features such as recommendations, content generation, analytics, and automation.                                                                        | AI functionality becomes part of the product experience. Inferencing services provide scalability, predictable performance, and faster deployment without requiring SaaS vendors to build AI infrastructure expertise. | Product Management, CTO, Engineering Leadership                               |
| Real-Time Inference APIs         | APIs that provide AI predictions, recommendations, classifications, fraud detection, or content generation to applications and services.                                                          | Requires high throughput, autoscaling, and consistent response times. Inferencing platforms optimize GPU utilization and API management for production workloads.                                                      | CTO, Application Development Teams, Platform Engineering                      |
| Code Assistants                  | AI tools that help developers generate code, explain code, perform reviews, create tests, and accelerate software delivery.                                                                       | Development activity can be highly variable. Inferencing as a Service enables organizations to provide enterprise-grade coding assistance without maintaining dedicated AI infrastructure.                             | VP Engineering, Software Development Leaders, Platform Engineering            |
| Internal Enterprise AI Platforms | Centralized AI platforms that provide reusable AI services, models, APIs, and governance capabilities across the organization.                                                                    | These platforms often support multiple business units and hundreds or thousands of users. Inferencing as a Service provides centralized operations, governance, monitoring, and cost management.                       | CIO, CTO, Chief Data Officer, AI Center of Excellence                         |

### Common Characteristics of Ideal Inferencing Workloads

The strongest candidates for Inferencing as a Service typically require:

- Real-time or near real-time responses
- Large Language Model (LLM) execution
- Variable or unpredictable demand
- High GPU consumption
- Enterprise security and compliance
- API-based integration into business applications
- Production-scale availability and reliability
- Ongoing model updates and lifecycle management



### Key Differentiators

- Enterprise-grade managed operations
- Flexible deployment models
- Support for open-source models
- Dedicated or shared inference capacity
- NVIDIA and AMD GPU flexibility
- OpenAI-compatible integration model
- Kubernetes-native scalability
- Private infrastructure support

### Resource Links

[Solution Hub](#)

### Contacts

**Offering Managers:**

Amine Badaoui: [Amine.Badaoui@rackspace.com](mailto:Amine.Badaoui@rackspace.com)

**Product Marketing:**

Sharon Varalli: [Sharon.Varalli@rackspace.com](mailto:Sharon.Varalli@rackspace.com)

# Rackspace Inferencing-as-a-Service

Accelerate Enterprise AI Adoption Without Increasing Operational Complexity

## Discovery Questions

| Use Case                         | Description                                                                                                                                                                                       | Why it's a good candidate for IaaS                                                                                                                                                                                     | Primary owner                                                                 |
|----------------------------------|---------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|-------------------------------------------------------------------------------|
| AI Copilots                      | AI assistants that help employees perform tasks, answer questions, generate content, summarize information, and automate workflows. Examples include HR, finance, sales, and healthcare copilots. | Requires low-latency responses, scalable GPU resources, and continuous access to LLMs. Usage patterns can fluctuate significantly, making elastic inferencing highly cost-effective.                                   | Chief Digital Officer, CIO, Business Unit Leaders, AI/Innovation Teams        |
| Enterprise Chatbots              | Internal or external conversational AI systems that answer employee, customer, or partner questions using enterprise knowledge.                                                                   | Demand varies widely throughout the day. Managed inferencing provides consistent performance, high availability, and simplified model operations without requiring dedicated AI infrastructure teams.                  | CIO, Customer Experience Leaders, IT Operations, Digital Transformation Teams |
| Customer Service Automation      | AI-powered virtual agents handling customer inquiries, ticket routing, call summarization, and self-service interactions.                                                                         | High transaction volumes require scalable, production-grade inferencing with predictable latency and uptime. This allows organizations to scale during peak customer demand.                                           | VP Customer Service, Contact Center Director, COO, CIO                        |
| Document Summarization           | AI-generated summaries of contracts, claims, medical records, financial reports, research documents, and other large content repositories..                                                       | Processing often occurs in bursts and requires substantial compute resources. Inferencing as a Service enables rapid scaling without overprovisioning GPU infrastructure.                                              | Legal Operations, Healthcare Informatics, Knowledge Management Teams, CIO     |
| Enterprise Search and RAG        | Retrieval-Augmented Generation systems that combine enterprise search with LLM-generated responses using company-specific data.                                                                   | Requires constant inferencing requests against vector databases and enterprise content repositories. It ensures low-latency query performance and simplifies operational complexity.                                   | CIO, Chief Data Officer, Knowledge Management Teams, Enterprise Architecture  |
| AI-Powered SaaS Applications     | Commercial software products embedding AI features such as recommendations, content generation, analytics, and automation.                                                                        | AI functionality becomes part of the product experience. Inferencing services provide scalability, predictable performance, and faster deployment without requiring SaaS vendors to build AI infrastructure expertise. | Product Management, CTO, Engineering Leadership                               |
| Real-Time Inference APIs         | APIs that provide AI predictions, recommendations, classifications, fraud detection, or content generation to applications and services.                                                          | Requires high throughput, autoscaling, and consistent response times. Inferencing platforms optimize GPU utilization and API management for production workloads.                                                      | CTO, Application Development Teams, Platform Engineering                      |
| Code Assistants                  | AI tools that help developers generate code, explain code, perform reviews, create tests, and accelerate software delivery.                                                                       | Development activity can be highly variable. Inferencing as a Service enables organizations to provide enterprise-grade coding assistance without maintaining dedicated AI infrastructure.                             | VP Engineering, Software Development Leaders, Platform Engineering            |
| Internal Enterprise AI Platforms | Centralized AI platforms that provide reusable AI services, models, APIs, and governance capabilities across the organization.                                                                    | These platforms often support multiple business units and hundreds or thousands of users. Inferencing as a Service provides centralized operations, governance, monitoring, and cost management.                       | CIO, CTO, Chief Data Officer, AI Center of Excellence                         |

**Contacts**

**Offering Managers:**

Amine Badaoui: [Amine.Badaoui@rackspace.com](mailto:Amine.Badaoui@rackspace.com)

**Product Marketing:**

Sharon Varalli: [Sharon.Varalli@rackspace.com](mailto:Sharon.Varalli@rackspace.com)

